

# Mercari Price Suggestion Challenge

*Ali Ezzat, Mohammed Ali*

*December 30, 2017*

## Contents

<b>About</b>	<b>1</b>
<b>Exploratory Data Analysis</b>	<b>1</b>
Data overview . . . . .	1
Price (The response/target variable) . . . . .	2
Item Condition . . . . .	4
Shipping . . . . .	6
Item Categories . . . . .	8
<b>Model Building</b>	<b>17</b>
<b>Conclusion</b>	<b>17</b>
<b>Credit</b>	<b>17</b>
<b>eBay acronyms</b>	<b>17</b>

## About

Mercari's challenge is to build an algorithm that automatically suggests the right product prices. You'll be provided user-inputted text descriptions of their products, including details like product category name, brand name, and item condition.

## Exploratory Data Analysis

### Data overview

- Load training data

```
train <- fread("Mercari-Price-Suggestion-Challenge/data/train.tsv", sep = "\t",
stringsAsFactors = FALSE, showProgress = FALSE)
```

- Inspect structure

```
glimpse(train)
```

```
## Observations: 1,482,535
## Variables: 8
## $ train_id      <int> 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13...
## $ name          <chr> "MLB Cincinnati Reds T Shirt Size XL", "Raze...
## $ item_condition_id <int> 3, 3, 1, 1, 1, 3, 3, 3, 3, 2, 1, 2, 1, 3, ...
## $ category_name   <chr> "Men/Tops/T-shirts", "Electronics/Computers ...
## $ brand_name     <chr> "", "Razer", "Target", "", "", "", "Acacia S...
## $ price          <dbl> 10, 52, 10, 35, 44, 59, 64, 6, 19, 8, 8, 34, ...
```

```
## $ shipping      <int> 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, ...
## $ item_description <chr> "No description yet", "This keyboard is in g..."
```

At the first glance, excluding the *train\_id* column, the numeric data are only 3 columns one of them is the response variable *price*, while the other two are factor data which mean that we have a very limited number of features and that mean we will need the help of the other four text predictors to build our model. So let us investigate these features one by one to see what we can do.

## Price (The response/target variable)

Let's start with an analysis of the response0 variable: price. First, the range of item prices:

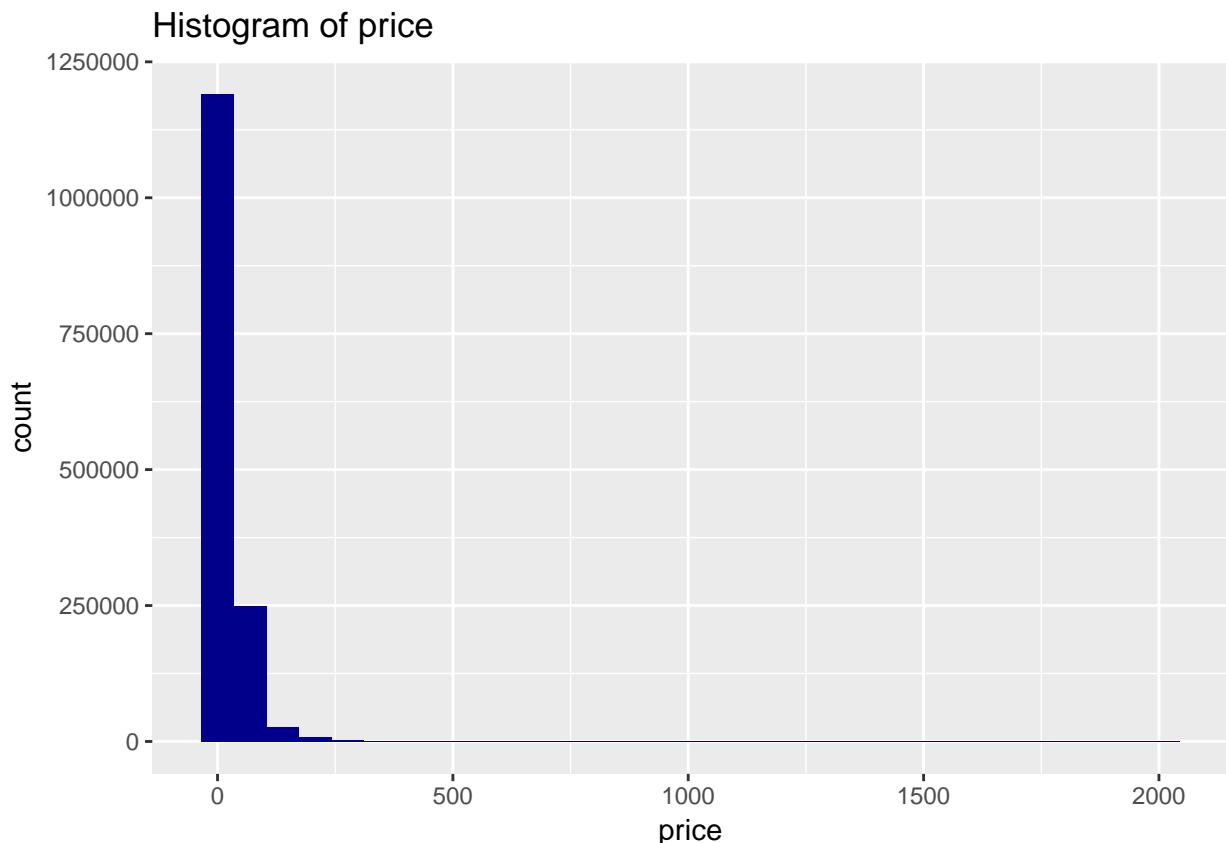
```
summary(train$price)
```

```
##   Min. 1st Qu. Median Mean 3rd Qu. Max.
##   0.00 10.00 17.00 26.74 29.00 2009.00
```

It seems there are items are given as gifts (min. is 0), and we have a few very expensive items as well. Let us visualize it for more clear picture

```
ggplot(data = train, aes(x = price)) +
  geom_histogram(fill = 'darkblue') +
  labs(title = 'Histogram of price')
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



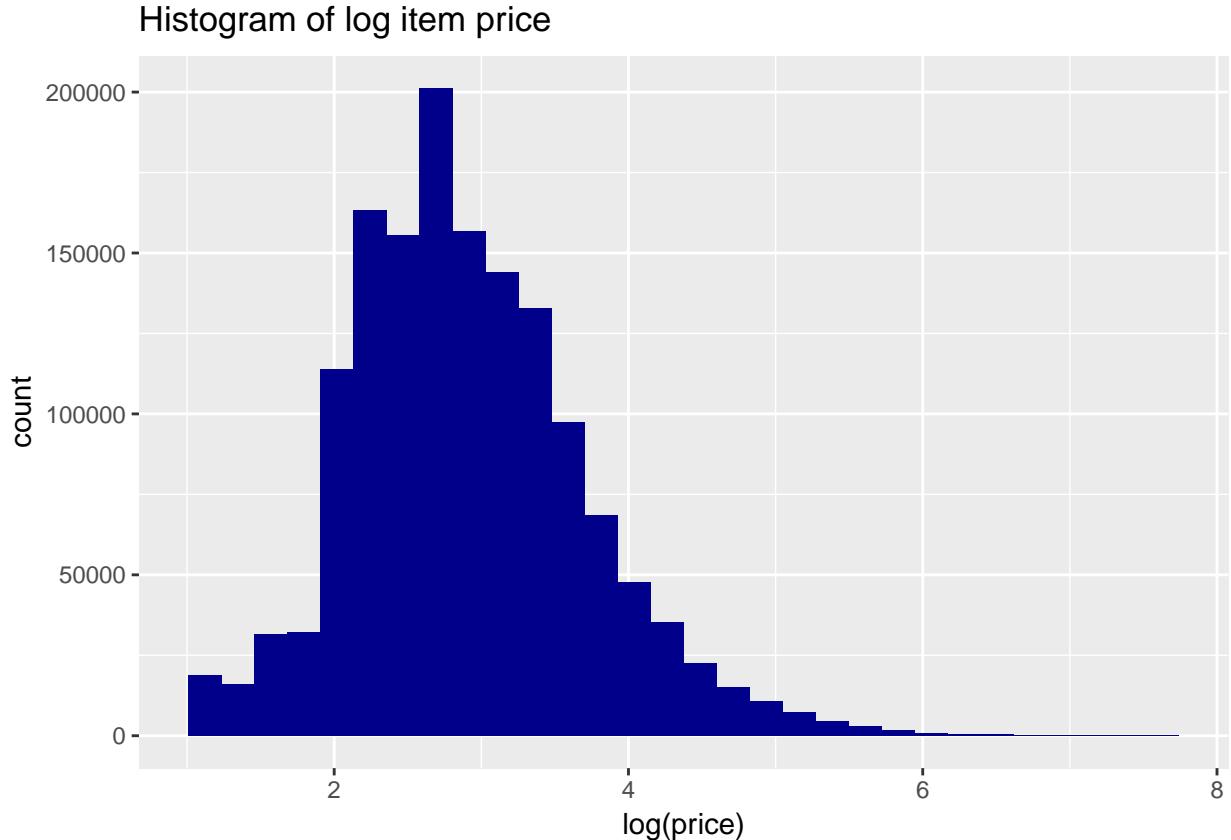
It seems because we have a very expensive items (and very view) we will need to take the log for better analysis

```

ggplot(data = train, aes(x = log(price))) +
  geom_histogram(fill = 'darkblue') +
  labs(title = 'Histogram of log item price')

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Removed 874 rows containing non-finite values (stat_bin).

```



Now, it is much better and clear, but it seems taking log of prices made the free/gift items to be omitted, as taking log for 0 is undefined, so we will have to add a dummy 1 to include it with us in the plot, consider it a tax :P .

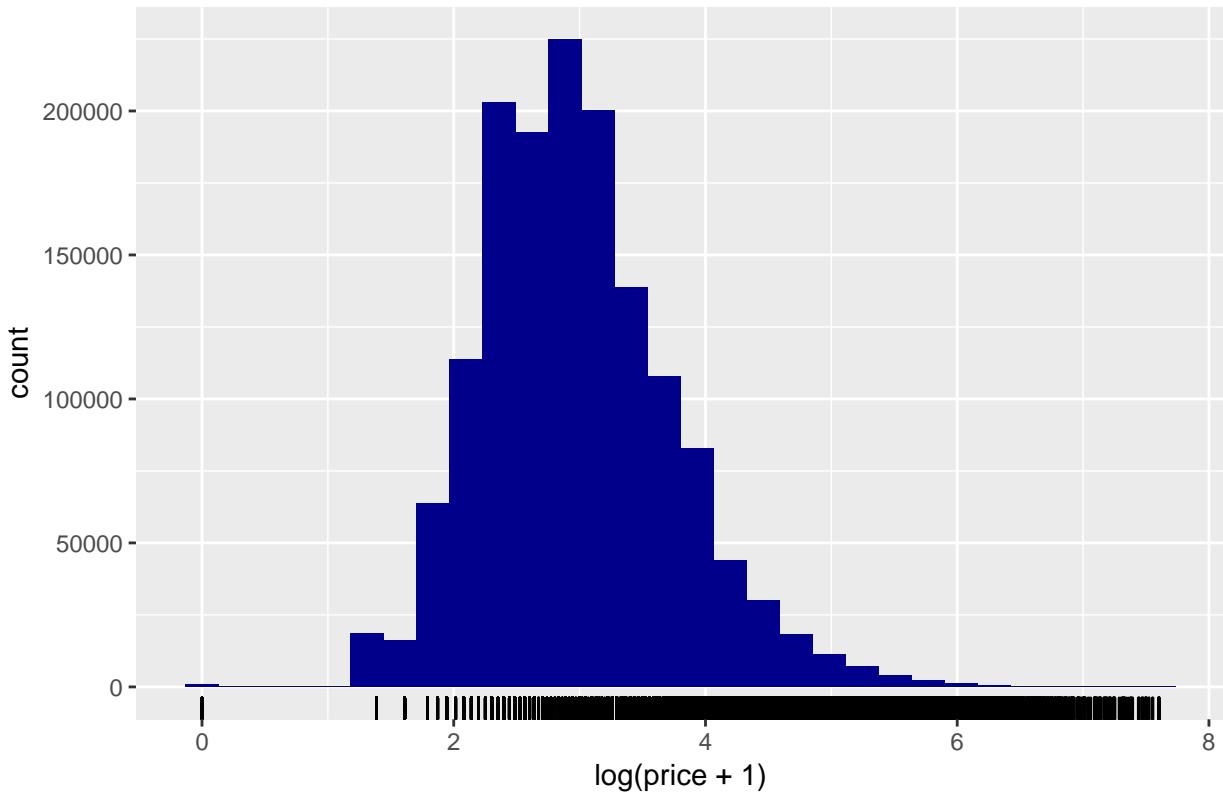
```

ggplot(data = train, aes(x = log(price + 1))) +
  geom_histogram(fill = 'darkblue') +
  labs(title = 'Histogram of log item price + 1') +
  geom_rug()

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

```

Histogram of log item price + 1



So, finally we have the gift/free items in the plot and it seems like an outlier along with the pricy data, as we have a nearly normal distribution for the prices. Let us investigate the distribution more. It seems the data is centered between 3 and 7, as also the following statistics confirm.

```
summary(log(train$price + 1))

##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
## 0.000   2.398   2.890   2.979   3.401   7.606
```

## Item Condition

Item condition is a factor data, so let us convert it to a factor first.

```
#item condition factor
train$item_condition_id <- as.factor(train$item_condition_id)
levels(train$item_condition_id)

## [1] "1" "2" "3" "4" "5"
```

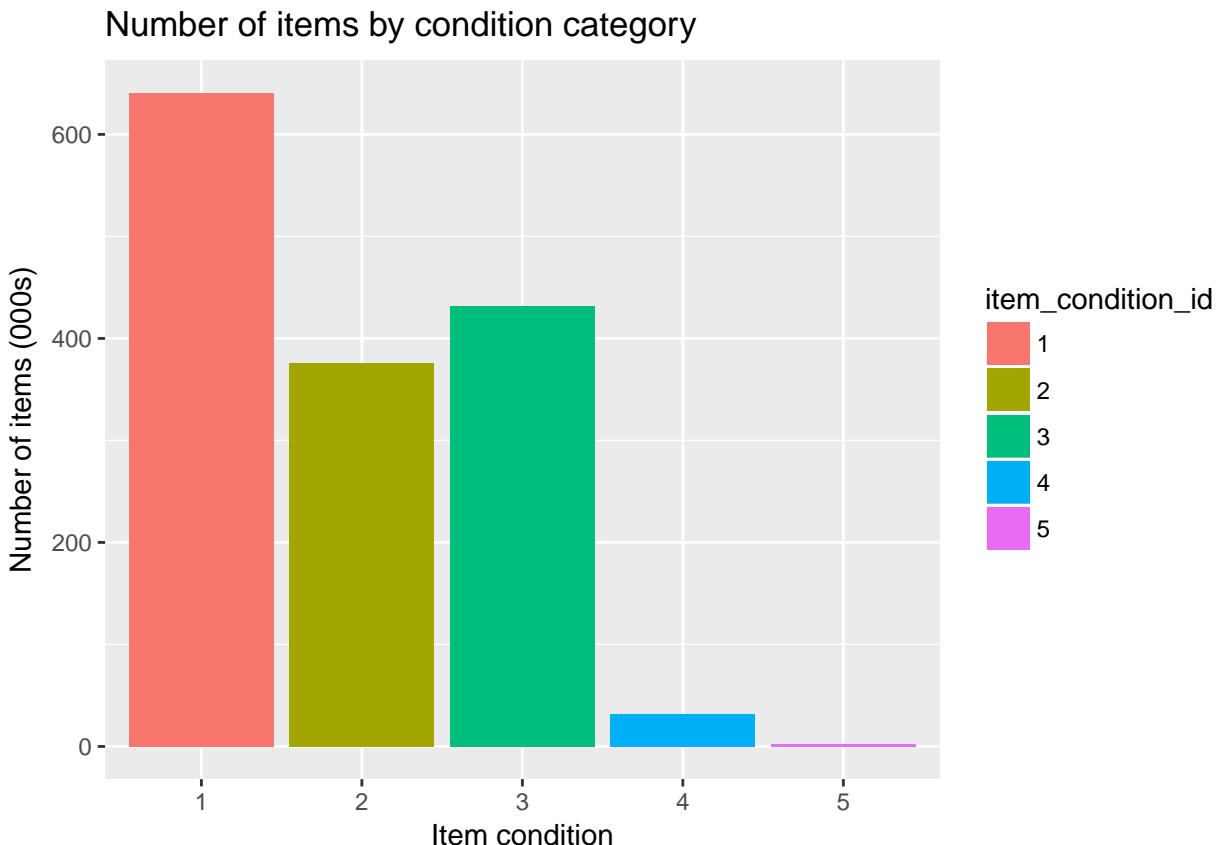
Now, let us see the statistics summary of items conditions

```
table(train$item_condition_id)

##
##      1      2      3      4      5
## 640549 375479 432161 31962   2384
```

We can confirm it more by the following plot

```
train[, .N, by = item_condition_id] %>% ggplot(aes(x = item_condition_id, y = N/1000,
  fill = item_condition_id)) + geom_bar(stat = "identity") + labs(x = "Item condition",
  y = "Number of items (000s)", title = "Number of items by condition category")
```



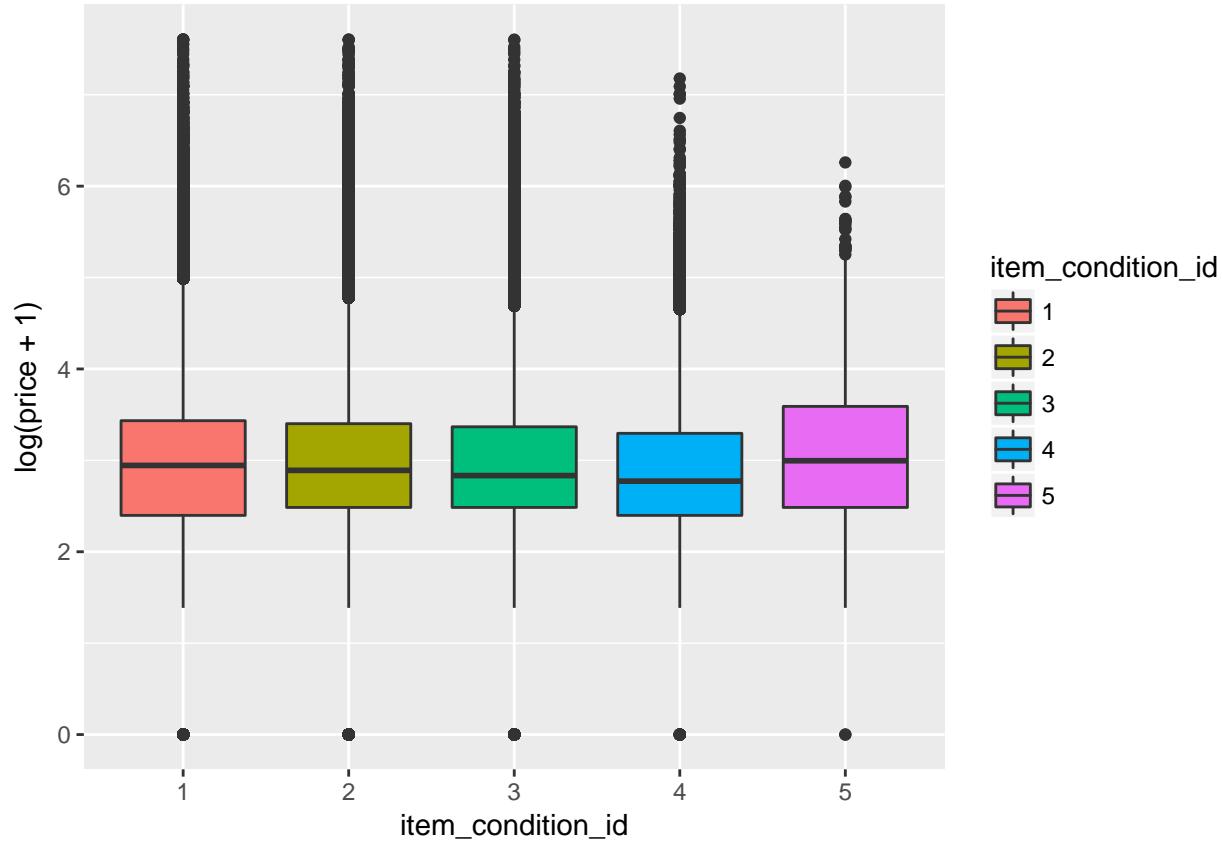
It seems most of the items are in condition 1 which is, I do not know there is no ordinal description in the competition. So we do not know if it 1 is the best or the worst. However, thanks to kaggler @Juraj for pointing out that in fact condition 1 is the best and 5 is the worst.

Now, let us compare the *item conditions* predictor against the response variable *price*

```
train[, .(N, median_price = median(price)), by = item_condition_id][order(item_condition_id)]
```

item_condition_id	N	median_price
1	640549	18
2	375479	17
3	432161	16
4	31962	15
5	2384	19

```
ggplot(data = train, aes(x = item_condition_id, y = log(price + 1), fill = item_condition_id)) +
  geom_boxplot()
```



It seems that item condition is not the main contributor to the price as the best price at condition 5 with few items and the second best price at condition 1 with a lot of items.

## Shipping

It is the second numeric facotr predictor. It is simply 1 when shipping item is paied by the seller and 0 otherwise. Let us convert it to a facotr and see its statistics

```
#shpping
train$shipping <- as.factor(train$shipping)
levels(train$shipping)

## [1] "0" "1"

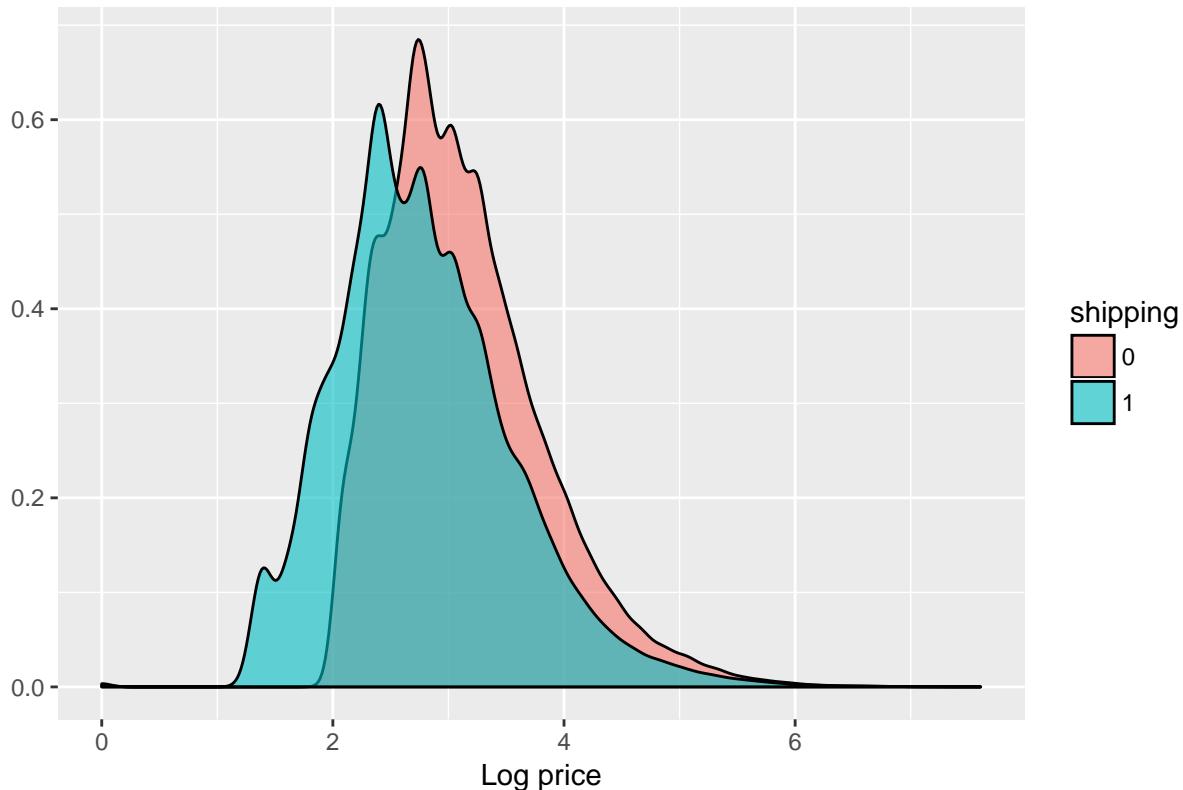
table(train$shipping)

##
##      0      1
## 819435 663100
```

It does not seem contributing that much to the reponse variable, so let us investigate its relation with the *price*

```
train %>%
  ggplot(aes(x = log(price+1), fill = shipping)) +
  geom_density(adjust = 2, alpha = 0.6) +
  labs(x = 'Log price', y = '', title = 'log(price) vs. shipping')
```

## log(price) vs. shipping



It seems that if you are going to pay for the shipping you will have a little lower price, but not that much right !! # Brand Clearly the two numeric predictor factors are not contributing so much into the response variable. How many brands we have?

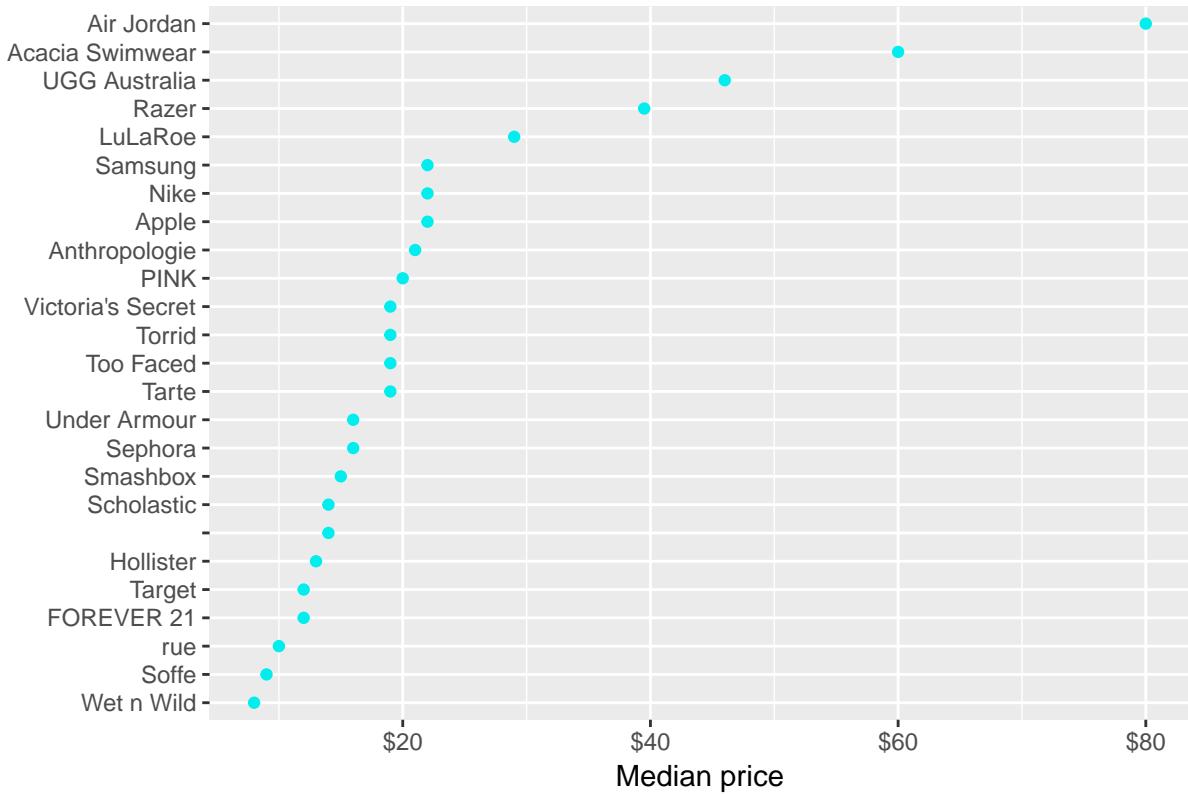
```
length(unique(train$brand_name))
```

```
## [1] 4810
```

Wow, that is a lot. So, let us try to investigate the text predictors and begin with the *brand* predictor.

```
train[, .(median_price = median(price)), by = brand_name] %>%
  head(25) %>%
  ggplot(aes(x = reorder(brand_name, median_price), y = median_price)) +
  geom_point(color = 'cyan2') +
  scale_y_continuous(labels = scales::dollar) +
  coord_flip() +
  labs(x = '', y = 'Median price', title = 'Top 25 most expensive brands')
```

## Top 25 most expensive brands



OK, finally we have a strongly contributer predictor. It seems that brands affect prices as it should be.

## Item Categories

Now it is *item categories* turn, we expect a lot from this variable. Let us begin by see how many categoris are there?

```
length(unique(train$category_name))
```

```
## [1] 1288
```

OK, there are a lot of categories here. So, how it looks like?

```
sort(table(train$category_name), decreasing = TRUE)[1:10]
```

```
##
##          Women/Athletic Apparel/Pants, Tights, Leggings      60177
##          Women/Tops & Blouses/T-Shirts                      46380
##          Beauty/Makeup/Face                                34335
##          Beauty/Makeup/Lips                                29910
##          Electronics/Video Games & Consoles/Games        26557
##          Beauty/Makeup/Eyes                                25215
```

```

## Electronics/Cell Phones & Accessories/Cases, Covers & Skins
##                                         24676
## Women/Underwear/Bras
##                                         21274
## Women/Tops & Blouses/Blouse
##                                         20284
## Women/Tops & Blouses/Tank, Cami
##                                         20284

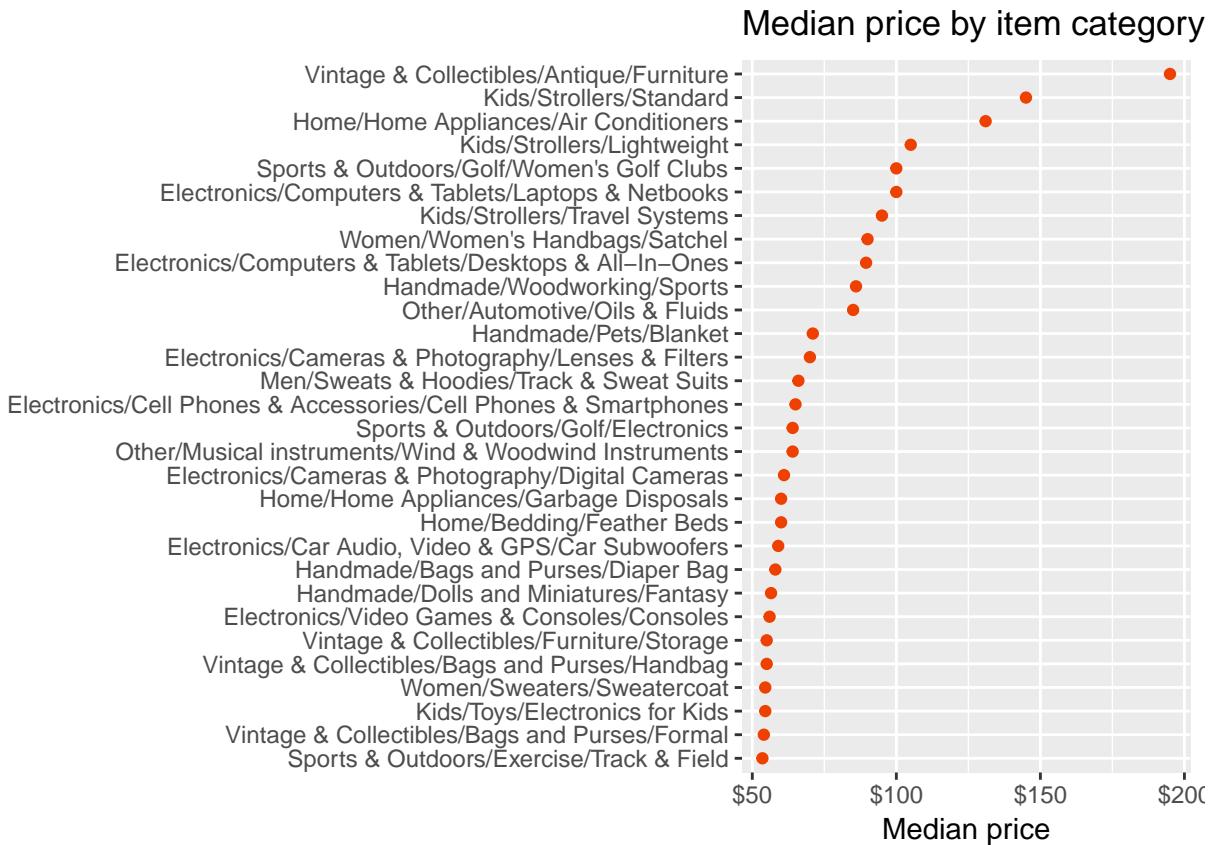
```

OK, there are a lot for women here :). It seems that we have a 3 level of category here, so let us investigate more. Now let us see how categories influence *price*

```

train[, .(median = median(price)), by = category_name][order(median, decreasing = TRUE)][1:30] %>%
  ggplot(aes(x = reorder(category_name, median), y = median)) + geom_point(color = "orangeRed2") +
  coord_flip() + labs(x = "", y = "Median price", title = "Median price by item category (Top 30)") +
  scale_y_continuous(labels = scales::dollar)

```



As expected, categories are contributing so much into the response variable *price* with *Vintage & Collectibles* in the top followed by *kids toys* !! So let us split each category to its level to see how each level contributes in the *price*

```

train[, `:=`(.data.table`(.train, , `:=`(.c("level_1_cat", "level_2_cat", "level_3_cat"), tstrsplit(train$category_name,
  split = "/")))]
```

```

## Warning in `.[.data.table`(.train, , `:=`(.c("level_1_cat", "level_2_cat", :
## Supplied 3 columns to be assigned a list (length 5) of values (2 unused)

```

```

head(train[, c("level_1_cat", "level_2_cat", "level_3_cat")])

##      level_1_cat      level_2_cat      level_3_cat
## 1:       Men          Tops        T-shirts
## 2: Electronics Computers & Tablets Components & Parts
## 3:     Women      Tops & Blouses        Blouse
## 4:     Home      Home Décor Home Décor Accents
## 5:     Women       Jewelry      Necklaces
## 6:     Women         Other        Other

```

We can now convert them to factors for easier and better analysis

```

train$level_1_cat <- as.factor(train$level_1_cat)
train$level_2_cat <- as.factor(train$level_2_cat)
train$level_3_cat <- as.factor(train$level_3_cat)

```

Now let us investigate each level #### level 1 The statistics of level1

```

length(levels(train$level_1_cat))

## [1] 10

table(train$level_1_cat)

## 
##      Beauty      Electronics      Handmade
##      207828          122690          30842
##      Home          Kids            Men
##      67871           171689          93680
##      Other      Sports & Outdoors Vintage & Collectibles
##      45351             25342          46530
##      Women
##      664385

```

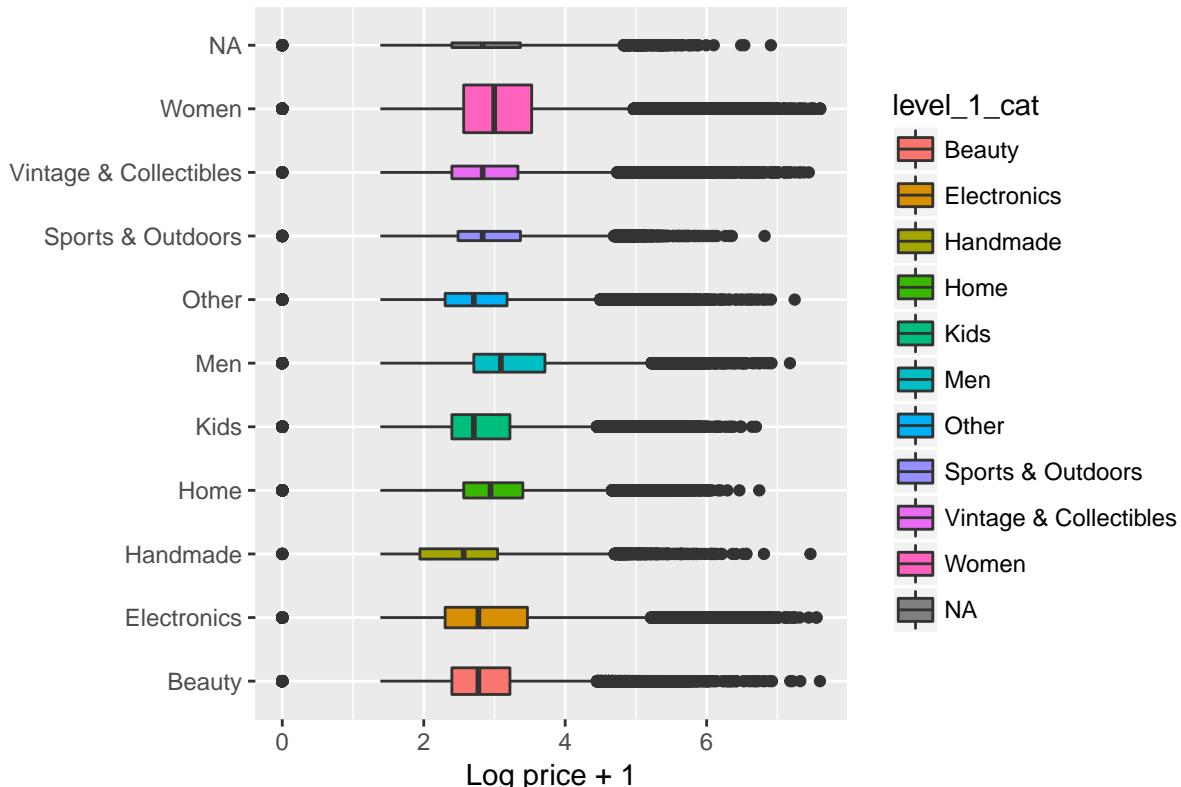
So, we have 10 items, let us its relation with *price*

```

train %>%
  ggplot(aes(x = level_1_cat, y = log(price+1), fill = level_1_cat)) +
  geom_boxplot(varwidth = TRUE) +
  coord_flip() +
  labs(x = '', y = 'Log price + 1', title = 'Boxplot of price by top-level category')

```

Boxplot of price by top-level category



So, we have items with no category but its prices are higher than **Homemade** items which is the *lowest* and less than **Men** items which are surprisingly higher than **Women** items in average. However, **Women** items still alot more than any other item.

## level 2

Statistics of level 2

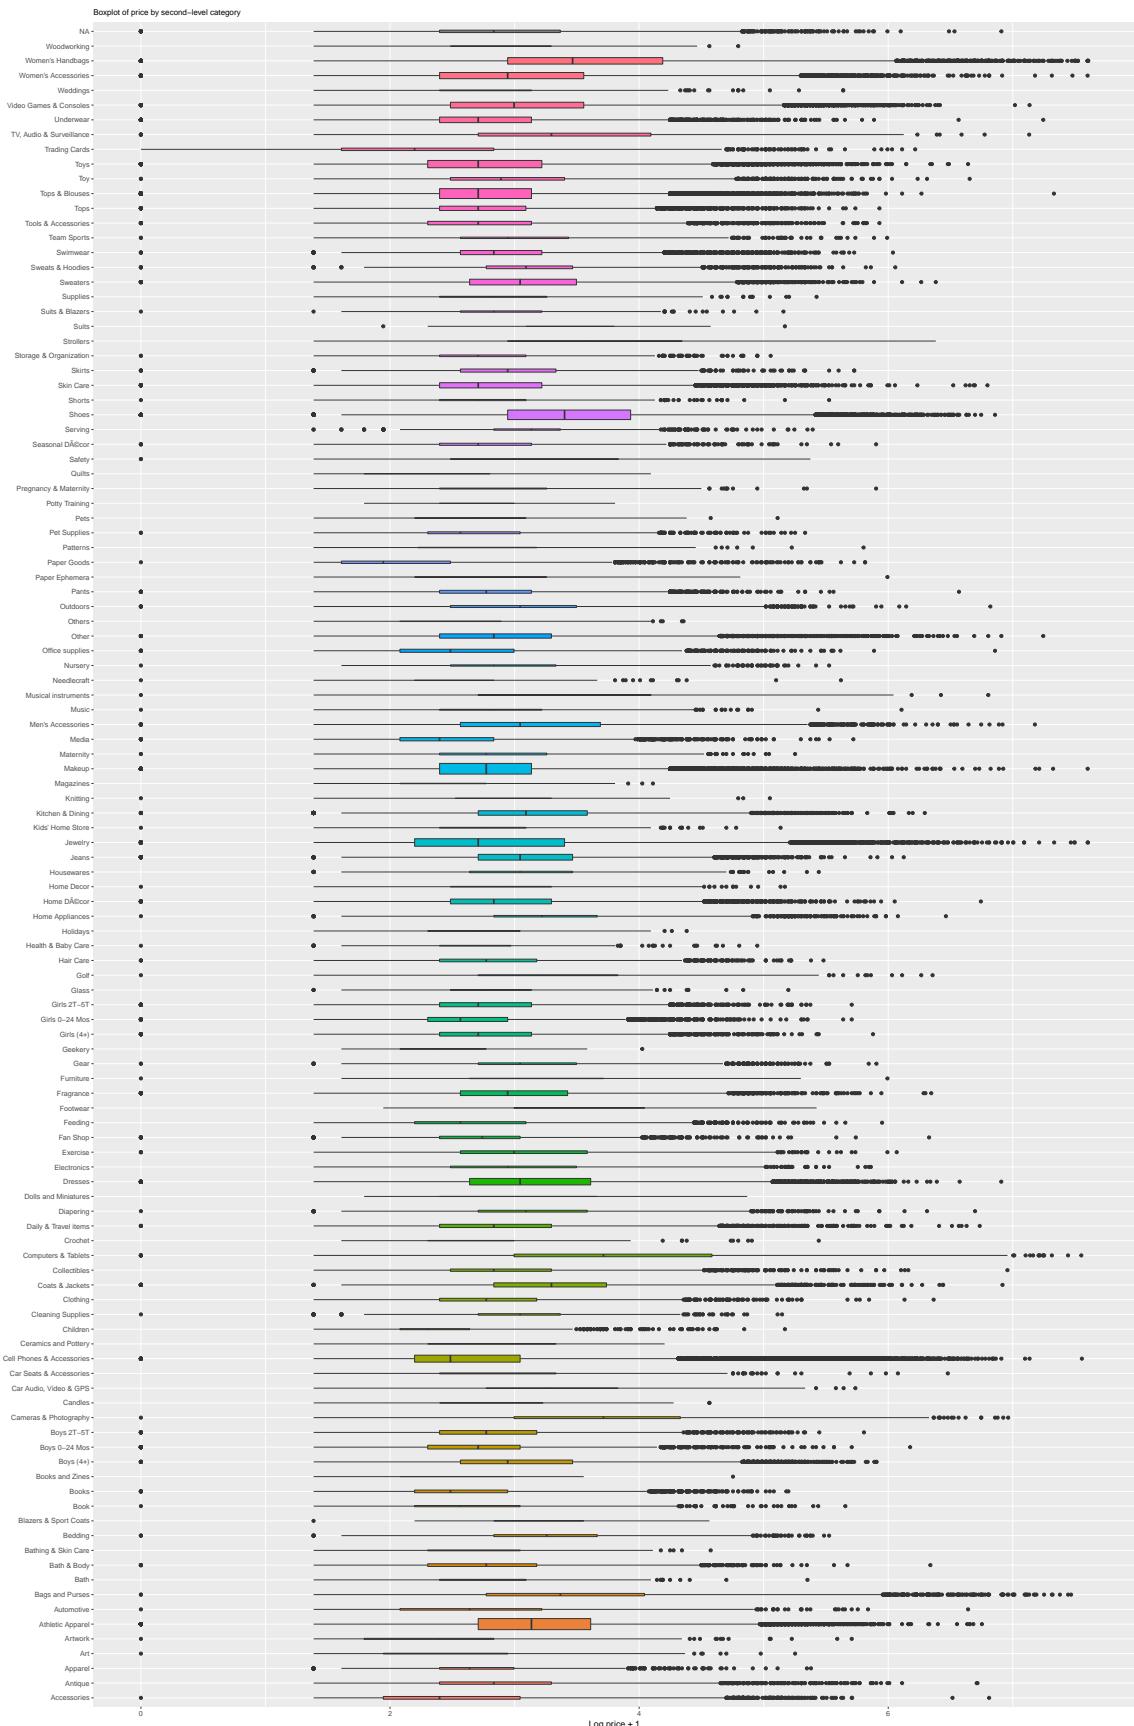
```
length(levels(train$level_2_cat))
## [1] 113
table(train$level_2_cat)[1:10]

##
##      Accessories          Antique        Apparel          Art
##             8213            6093            2918            656
##      Artwork  Athletic Apparel    Automotive Bags and Purses
##             1264           134383            2480            6338
##      Bath      Bath & Body
##             1192            7758
```

We have a lot more subcategories in level 2, let us see their prices.

```
train %>%
  ggplot(aes(x = level_2_cat, y = log(price+1), fill = level_2_cat)) +
  geom_boxplot(varwidth = TRUE) +
  coord_flip()
```

```
labs(x = '', y = 'Log price + 1', title = 'Boxplot of price by second-level category') +  
theme(legend.position="none", plot.title = element_text(size=10))
```

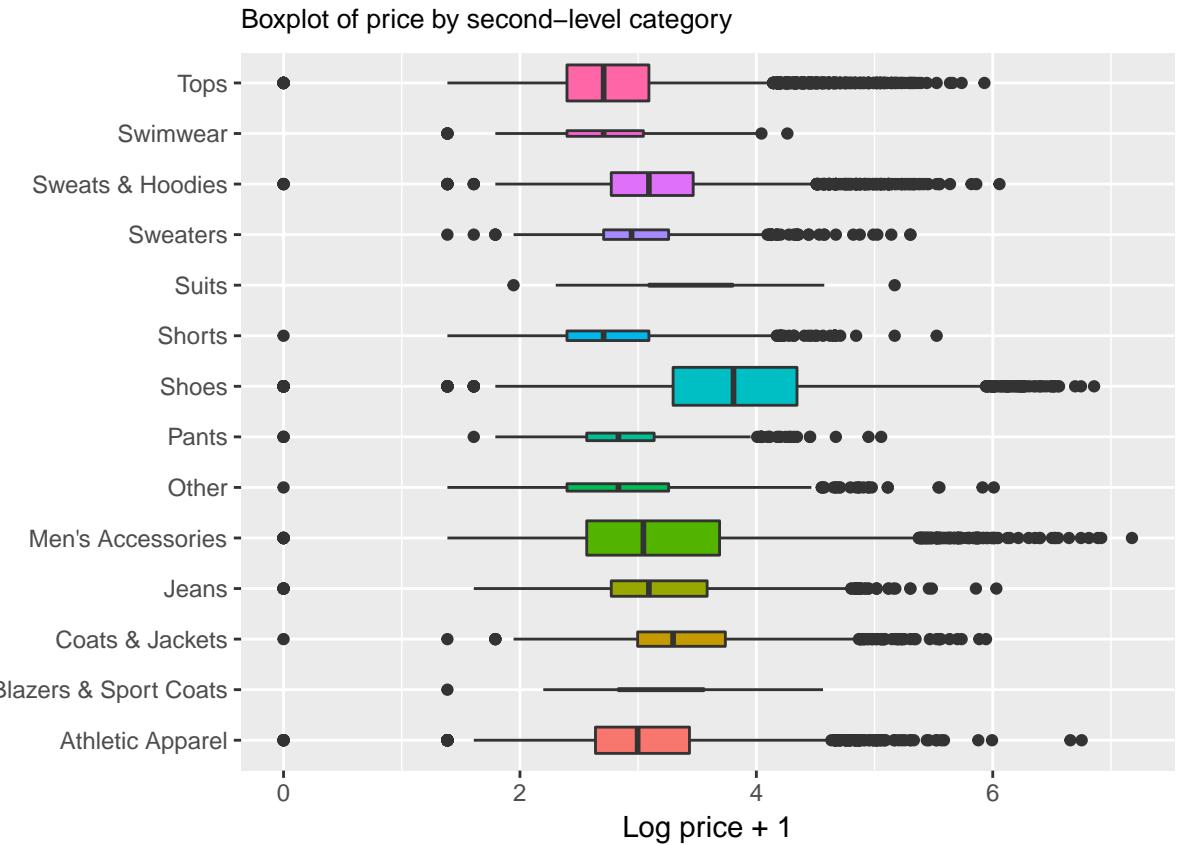


Still

*Women* and *Kids* items are at the top. We can also see second level category items distribution against *price* for specific first level, let us compare between *Men* and *Women* items

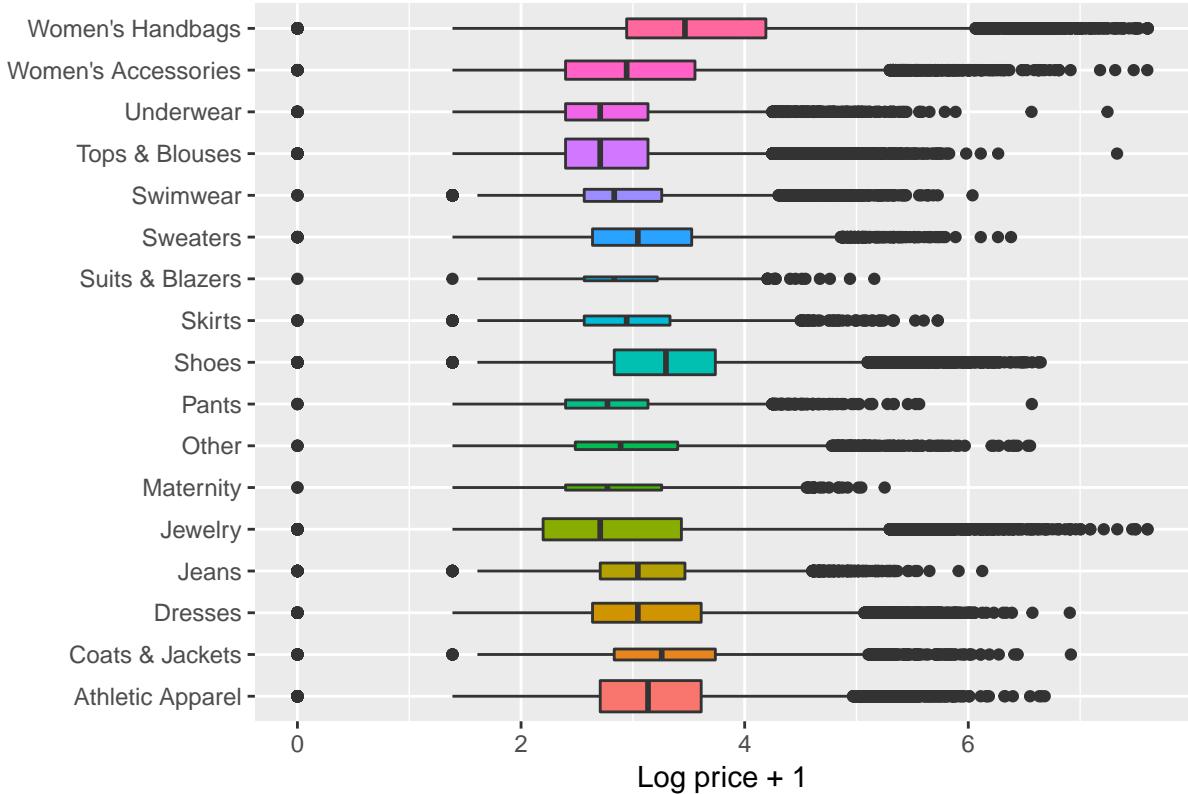
```
train %>%
  filter(level_1_cat == "Men") %>%
  ggplot(aes(x = level_2_cat, y = log(price+1), fill = level_2_cat)) +
  geom_boxplot(varwidth = TRUE) +
  coord_flip() +
  labs(x = '', y = 'Log price + 1', title = 'Boxplot of price by second-level category') +
  theme(legend.position="none", plot.title = element_text(size=10))
```

## Warning: package 'bindrcpp' was built under R version 3.4.2



```
train %>%
  filter(level_1_cat == "Women") %>%
  ggplot(aes(x = level_2_cat, y = log(price+1), fill = level_2_cat)) +
  geom_boxplot(varwidth = TRUE) +
  coord_flip() +
  labs(x = '', y = 'Log price + 1', title = 'Boxplot of price by second-level category') +
  theme(legend.position="none", plot.title = element_text(size=10))
```

Boxplot of price by second-level category



It seems Men's shoes and accessories are what *Men* items are higher than *Women* items. Otherwise, *Women* items are more pricy in general.

### level3

statistics of level3

```
length(levels(train$level_3_cat))

## [1] 870

table(train$level_3_cat)[1:10]

##          100 Years or Older      50 To 75 Years      75 To 100 Years
##                      140                      92                      42
##          A-Line      Above Knee, Mini      Accessories
##                      756                  20082                  10297
##          Accessory      Aceo      Action Figure
##                      31                      1                  5354
## Action Figures & Statues      12524
```

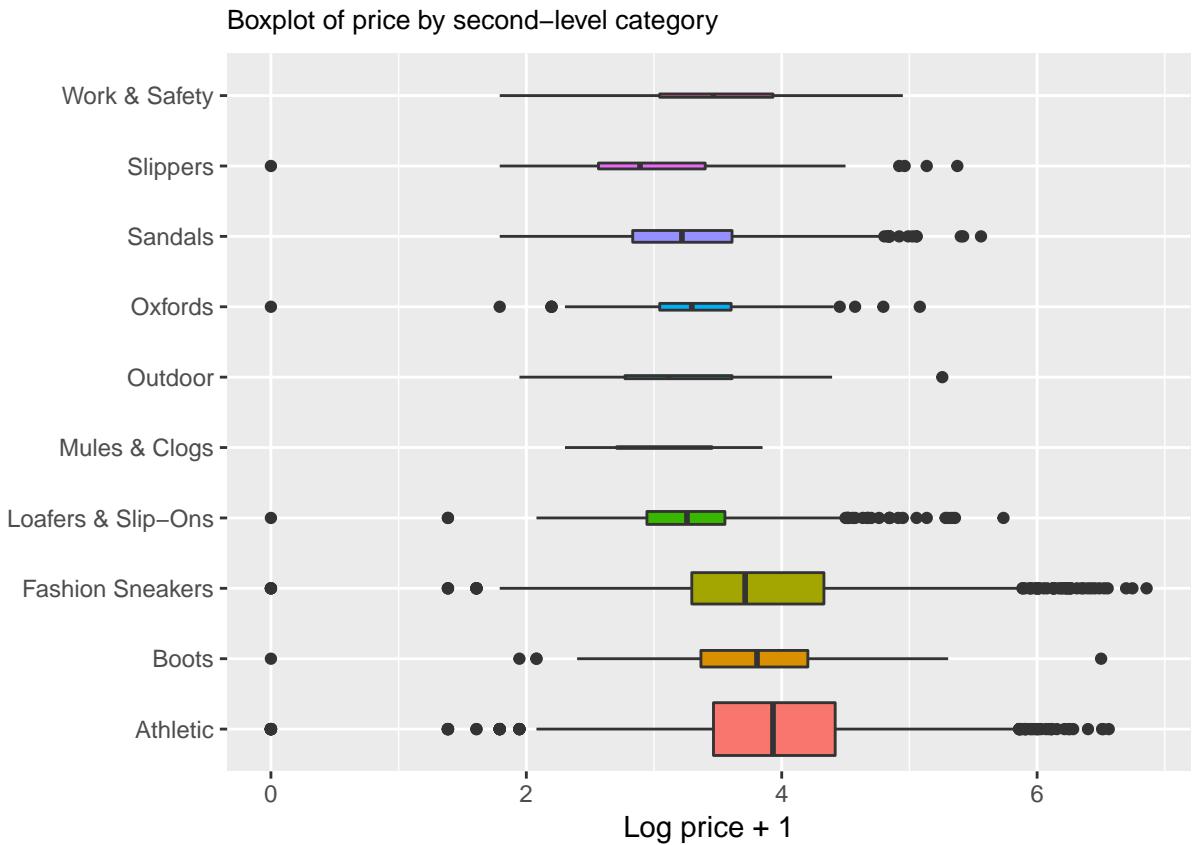
It will hard to visualize all these data at once, so let us seen what *brands* are pricy for *Men* and *Women* pricy items

```
train %>%
  filter(level_1_cat == "Men" & level_2_cat == "Shoes") %>%
```

```

ggplot(aes(x = level_3_cat, y = log(price+1), fill = level_3_cat)) +
  geom_boxplot(varwidth = TRUE) +
  coord_flip() +
  labs(x = '', y = 'Log price + 1', title = 'Boxplot of price by second-level category') +
  theme(legend.position="none", plot.title = element_text(size=10))

```



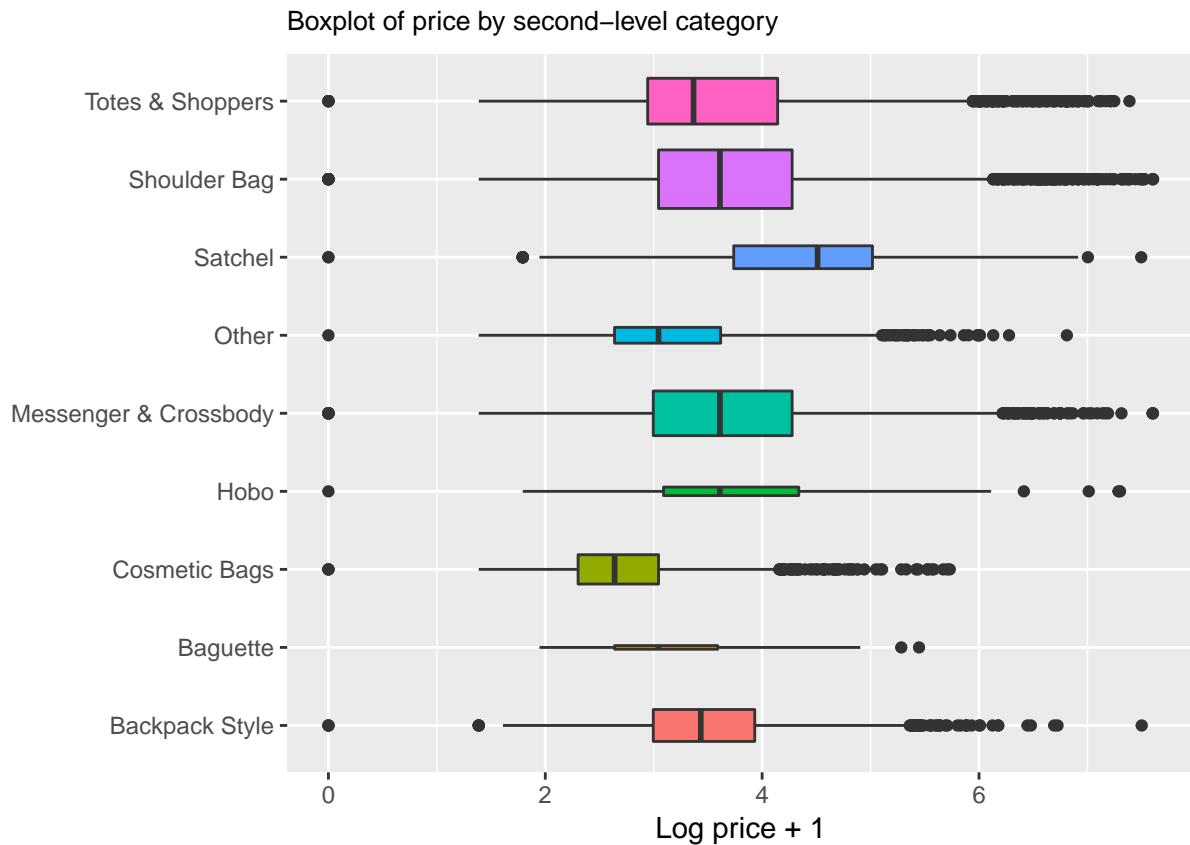
Ok, we have a lot of men are going to Gym here :).

and for women

```

train %>%
  filter(level_1_cat == "Women" & level_2_cat == "Women's Handbags") %>%
  ggplot(aes(x = level_3_cat, y = log(price+1), fill = level_3_cat)) +
  geom_boxplot(varwidth = TRUE) +
  coord_flip() +
  labs(x = '', y = 'Log price + 1', title = 'Boxplot of price by second-level category') +
  theme(legend.position="none", plot.title = element_text(size=10))

```



## Model Building

### Conclusion

### Credit

Troy Walters

### eBay acronyms

A-E B&W: black and white

BC: back cover (usually used as a description for books)

BIN: Buy It Now

CIP: customer initiated payment

DOA: dead on arrival (an item that doesn't work or is broken when it's received)

DSR: detailed seller rating (additional Feedback ratings buyers can give sellers)

EST: Eastern Standard Time

EUC: excellent used condition

F-I FAQ: frequently asked questions (a list of questions with answers.)

FB: Feedback

FC: fine condition

FOB: freight on board (usually means something has shipped)

FS: full screen (usually applied to a DVD or video format)

FVF: final value fee

G: good condition

GBP: Great Britain pounds

GU: gently used (item that has been used but shows little wear, accompanied by explanation of wear)

HP: home page

HTF: hard to find

HTML: hypertext markup language (the language used to create web pages)

IE: Internet Explorer

IM: instant messaging

INIT: initials

ISP: Internet service provider (a company that gives you access to the Internet)

J-M JPG: JPEG (preferred file format for pictures on eBay, pronounced “jay-peg”)

LTBX: letterbox (video format that recreates a widescreen image)

LTD: limited edition

MNT: mint or in perfect condition (a subjective term that doesn't necessarily mean new)

MIB: mint in box

MIJ: made in Japan

MIMB: mint in mint box

MIMP: mint in mint package

MIP: mint in package

MNB: mint no box

MOC: mint on card

MOMC: mint on mint card

MONMC: mint on near mint card

MWBT: mint with both tags

MWMT: mint with mint tags

N-P NARU: not a registered user (also a suspended user)

NBW: never been worn

NC: no cover

NIB: new in box

NM: near mint

NOS: new old stock

NR: no reserve price (for an auction-style listing)

NRFB: never removed from box

NWT: new with tags

NWOB: new without box

NWOT: new without original tags

OEM: original equipment manufacturer

OOP: out of print

PST: Pacific Standard Time

Q-Z RET: retired

SCR: scratch

S/O: sold out

Sig: signature